

Neural Networks & Activation Functions

Sattam Altwaim

sattam.tuuaim@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology



KAUST Academy & Mawhiba
AI Olympiad Training (Beginners)
Week 3, Day 18, Self-study

6 October 2026

1. Where This Fits
2. Learning Outcomes
3. Activation Functions in Depth
4. What a ReLU Network Computes
5. Choosing Width and Depth
6. Output Layers and Their Losses
7. Olympiad Practice
8. Check Yourself
9. Limitations and What Comes Next
10. Summary
11. References

- ▶ **In the lecture (Day 15):** a neuron is a line plus a bend, the bend is the activation function, a network stacks layers of neurons, and the chain rule carries the slope back through every layer. Day 17 added that deep sigmoid networks do not train without BatchNorm, and that dials such as dropout change what a unit sees.
- ▶ **This deck adds:** the slopes of the three activations with numbers and what happens to them across five layers; what a “dead” unit is and how to find one; what a ReLU network actually computes (a sum of hinges, drawn on real cars and as a node graph); how to choose width and depth from the training and validation errors, with the rules of thumb and where they come from; and the two output-layer mistakes that cost the most points.
- ▶ **How to use it:** every frame stands alone and every number was measured. Read a frame, type its code into a notebook, change one number, and compare with the frame. The self-check at the end tells you whether to reread.

Where the simple version breaks

- ▶ Day 15 said “ReLU inside, sigmoid or softmax at the output” and moved on. Day 17 then showed five sigmoid layers that never left a loss of 2.3. Why exactly, and would ReLU have saved them? The answer is a number: at the start of training the gradient reaching the first sigmoid layer is 4,900 times smaller than the one at the output layer; with ReLU the two are about the same size.
- ▶ “The network bends” is not a mechanism, and “use 128 units” is not a design. With one input you can draw every piece: 6 hidden units fit 3,598 used-car prices with MAE 39,017 against 41,724 for a line, and the picture shows which unit bends where. On 2,400 games, 4 units underfit, 64 fit and 1,024 memorise: the shape is read off two curves, not guessed.
- ▶ A network with a softmax inside it *trains*: the loss falls and the accuracy climbs. It is still wrong, and the gap only shows next to the right version: a training loss of 1.49 against 0.035 after eight epochs.

After this deck you can


read a network’s activations and output layer the way Day 16 taught you to read its loss curve: as a dashboard with a cause behind every symptom.

By the end of this self-study deck, you will be able to:

- ▶ **Compute** the slope of sigmoid, tanh and ReLU at a given z , and explain with numbers why slopes below 1 vanish through layers.
- ▶ **Diagnose** a dead ReLU unit and a saturated sigmoid from the numbers a network prints, and say what to change.
- ▶ **Draw** what a one-input ReLU network computes: a sum of hinges, one kink per hidden unit.
- ▶ **Choose** width and depth from the training and validation errors, using the rules of thumb and their sources, and count the dials of the choice.
- ▶ **Pair** every output layer with its loss, and spot the two classic mistakes: a softmax inside the model, and an activation on a regression output.

Activations: **Activation Functions in Depth**

The Picture, With Today's Questions On It



Each circle is a neuron (Day 15): it sums its arrows, then bends. Today's questions, by number:


1. **Inside a hidden node:** what does the bend do to the error signal on its way back? (this section)
2. **A whole hidden layer:** what function do its units add up to? (a sum of hinges)
3. **How many units, how many layers?** A design choice with a measured answer (choosing width and depth)
4. **After the last node:** nothing. The loss does the softmax (output layers)

Backpropagation multiplies slopes (Day 15). The slope of the activation is one factor in every layer, so its size decides how much of the error signal survives the layer.


Activation	Slope formula	$z = 0$	$z = 2$	$z = 5$	$z = -5$
sigmoid	$\sigma(z) (1 - \sigma(z))$	0.25	0.105	0.0066	0.0066
tanh	$1 - \tanh^2(z)$	1	0.071	0.0002	0.0002
ReLU	1 if $z > 0$, else 0	0	1	1	0

- ▶ Sigmoid's slope is **never above 0.25**, and it is tiny once $|z| > 4$: the unit has **saturated**. tanh reaches 1 at $z = 0$ but saturates just as hard.
- ▶ ReLU's slope is exactly 1 on the whole right side: the signal passes through unchanged, however large z is. On the left it is exactly 0: nothing passes.
- ▶ Check one by hand: $\sigma(2) = 0.881$, so the slope is $0.881 \cdot 0.119 = 0.105$. At $z = 5$: $\sigma(5) = 0.9933$, slope $0.9933 \cdot 0.0067 = 0.0066$.

Vanishing Gradients, By Hand and Measured



One activation slope per layer multiplies the error signal on its way back: through five sigmoid layers at best $0.25^5 = 0.00098$ of it reaches layer 1. With ReLU every factor is 1.



Source: 784-256-256-256-256-10 networks at initialisation, gradient of each layer's weights on one batch of 64 MNIST images; `scripts/day18_activations_figures.py`.

Measured: with sigmoid the first layer's gradient is 4,900 times smaller than the output layer's; with tanh, ReLU, or sigmoid plus BatchNorm the six layers stay within a factor of three. BatchNorm (Day 17) keeps every z near 0, where the sigmoid slope is at its 0.25 maximum.

```
z = torch.tensor([0., 2., 5., 50.])
s = torch.sigmoid(z)
print(s)           # 0.50 0.88 0.993 1.0
print(s * (1 - s)) # 0.25 0.105 0.0066 0
print(torch.tanh(z)) # 0.00 0.96 1.000 1.0
print(torch.relu(z)) # 0. 2. 5. 50.
```

At $z = 50$ the sigmoid's slope is exactly 0 in `float32`: the unit is frozen and no step will unfreeze it.


- ▶ Where do z values of 50 come from? From **unscaled inputs**: a pixel of 255 or a mileage of 200,000 times a weight of 0.1 is already far outside the live zone.
- ▶ That is the mechanism behind Day 15's "scale first": standardised inputs (mean 0, spread 1) keep the first layer's z within a few units of 0.
- ▶ For deeper layers, BatchNorm does the same job after every layer. For ReLU the live zone is the whole right side: one more reason it is the default.

Definition

A ReLU unit is **dead** when its $z = \mathbf{w} \cdot \mathbf{x} + b$ is negative for every input. Its output and its slope are always 0, so no gradient reaches its weights: it stays dead.

```
model.eval()
with torch.no_grad():
    z = model[0](X_val)
    on = (z > 0).float().mean(0)
dead = (on == 0).sum().item()
print(dead, "dead units of", len(on))
```

z is layer 1 before its ReLU; on is the share of rows on which each unit fires. A unit that fires on none is dead.



the grey unit: $z < 0$ for every row, so 0 on every arrow out and no gradient on any arrow in

- ▶ How units die: one big update (a large rate, an outlier batch) pushes b far negative; the unit never fires again.
- ▶ Measured on 784-128-10 after three epochs of SGD: rate 0.05 kills no unit (accuracy 0.944); 0.5 kills 21 of 128 (0.948); 1.0 kills 90 of 128 and the accuracy collapses to 0.546.
- ▶ The fix is upstream: a smaller rate, scaled inputs, BatchNorm. **Leaky ReLU** (`nn.LeakyReLU()`, slope 0.01 on the left) lets a unit come back.

Activations: **What a ReLU Network Computes**

With a single input x , hidden unit j computes $h_j = \text{ReLU}(w_j x + b_j)$ and the output adds $v_j h_j$. The function $v_j \text{ReLU}(w_j x + b_j)$ is a **hinge**: 0 on one side of the kink $x = -b_j/w_j$, a line of slope $v_j w_j$ on the other.


With numbers

$w = 1$, $b = -2$, $v = 3$: the hinge is $3 \cdot \text{ReLU}(x - 2)$. At $x = 0$: $\text{ReLU}(-2) = 0$, contribution 0. At $x = 2$: still 0 (the kink). At $x = 4$: $\text{ReLU}(2) = 2$, contribution 6. At $x = 5$: 9. Slope 3 to the right of 2, flat to the left.

With $w = -1$, $b = 2$, $v = 3$: $3 \cdot \text{ReLU}(2 - x)$, the mirror image: flat to the right of 2, slope -3 to the left.

- ▶ The output layer adds the hinges and a bias: $\hat{y} = b_2 + \sum_j v_j \text{ReLU}(w_j x + b_j)$. A sum of straight pieces joined at kinks is a **piecewise linear** curve: with m hidden units, at most m kinks.
- ▶ Training moves the kinks (through w_j, b_j) and the slopes (through v_j). A ReLU network never draws a true curve, only straight pieces, as many as it has units.


Six Hinges on 3,598 Cars



Source: 3,598 listings from 2015 on, at most 300,000 km, Kaggle turkibintalib/saudi-arabia-used-cars-dataset, 80/20 split; 1-6-1 ReLU network, L1 loss, Adam; scripts/day18_activations_figures.py.

- ▶ Left: the fit is straight pieces: a drop over the first 150,000 km, then nearly flat. It beats the line by 2,700 SAR of MAE (39,017 against 41,724) with no hand-made x^2 column: the bends are the units.
- ▶ Right: units 4 and 6 bend inside the range: a slope down, then zero past kinks near 153,000 and 283,000 km (dotted). Units 1, 3 and 5 kink outside the range, so here they are lines; unit 2 is almost zero. Their sum plus the output bias is the purple curve.

The Same Network as Nodes and Arrows



kinks in thousand km; "on below" means $w < 0$

- ▶ Each hidden node is one hinge: its kink is $-b_j/w_j$ (converted to km), the sign of w_j says on which side it is on, and the arrow to $\hat{y}$ carries its slope v_j .
- ▶ The two dark units kink inside the data, at 154 and 284 thousand km: they made the two bends of the previous frame. Units 1, 3 and 5 kink outside the range, so for every real car they are plain lines.
- ▶ Unit 2 is never on: for this data it is a dead unit, and its arrow carries nothing. Six dials bought, five used.
- ▶ The output node adds the contributions and b_2 . Reading a trained network is reading its dials; the graph is the map.

dials of a layer = inputs $\times$ units + units


dials of the network = the sum over its layers

Shape	Arithmetic	Dials	MNIST test accuracy
784-32-10	$784 \cdot 32 + 32 + 32 \cdot 10 + 10$	25,450	0.951
784-128-10	$784 \cdot 128 + 128 + 128 \cdot 10 + 10$	101,770	0.964
784-512-10	$784 \cdot 512 + 512 + 512 \cdot 10 + 10$	407,050	0.970
784-128-128-10	$101,770 - 1,290 + 128 \cdot 128 + 128 + 1,290$	118,282	0.963
784-128-128-128-10	$118,282 + 128 \cdot 128 + 128$	134,794	0.967

Accuracy after eight epochs of Adam at 0.001 on 20,000 MNIST training images, same seed for every shape; `scripts/day18_activations_figures.py`.

- The first layer owns almost everything: 784 inputs make each of its units expensive; a second layer of 128 costs 16,512 dials, a sixth of a 512-wide first layer. `sum(p.numel() for p in model.parameters())` counts them.
- Wider helped with diminishing returns (4 and 16 times the dials for 1.3 and 1.9 points); a second and third layer bought 0.963 and 0.967, no better than one wide layer. Why, and how to choose: the next section.

Activations: **Choosing Width and Depth**




Definition

The **capacity** of a network is how much it can bend: for ReLU networks, how many straight pieces it can make. Width adds pieces, depth lets layers fold each other's pieces, and every arrow above is one dial to fit.

- ▶ Too little capacity and the **training** error itself stays high: underfitting. Too much and the training error falls while the **validation** error does not follow: overfitting, Day 17's gap.
- ▶ One hidden layer can in principle bend into any shape if wide enough (Cybenko 1989; Hornik et al. 1989). The question is how many dials this data justifies, and the two errors answer it.


Measured: Three Regimes on 2,400 Games



Source: the Steam prices with Day 16's features, 2,400 / 600 split, one hidden layer of 4, 64 or 1,024 ReLU units, Adam at 0.001, L1 loss, 300 epochs; `scripts/day18_design_figures.py`.


- ▶ Left: 4 units cannot even fit the training games (training MAE 6.17, validation 6.53): underfitting, cured only by capacity. Middle: 64 units fit (5.49 and 6.42) and the validation curve stays flat for 300 epochs. Right: 1,024 units push the training MAE to 4.10 while the validation MAE climbs from 6.40 at epoch 31 to 6.92: memorising.
- ▶ Same data, recipe and epochs; only the width changed. The training curve alone would pick the biggest network; the validation curve picks the middle one.

Measured: the Sweep Over Width and Depth



Source: the same data and training recipe; widths 4 to 1,024 with one hidden layer, depths 1 to 6 with 64 units per layer; `scripts/day18.design_figures.py`.

- ▶ Width: the best validation MAE is flat from 16 units on (6.44 to 6.40) while the training MAE keeps falling (5.79 to 4.10): past 64 units the dials buy memorising. Depth at 64 units: the best score never moves (6.37 to 6.41), the final one climbs to 8.02 with six layers, and the best epoch drops from 178 to 11: deeper memorises faster.
- ▶ With early stopping even the 1,024-unit network scores 6.40: a too-big shape costs time and a risky late stop, not accuracy. Hence the rule “large enough, then regularise”, not “as small as possible”.




Source: the tent map $2 \text{ReLU}(x) - 4 \text{ReLU}(x - 0.5)$ applied once, twice and three times; `scripts/day18.design_figures.py`.

- ▶ One hidden layer of n units draws at most $n + 1$ straight pieces on a line (the hinge frames). A second layer bends that bent curve again: the pieces **multiply**.
- ▶ The tent map is one layer of 2 units ($w = 1, 1$; $b = 0, -0.5$; $v = 2, -4$). Stacked k times it folds $[0, 1]$ in half k times: 2^k pieces from $2k$ units, where one layer needs $2^k - 1$ units; for $k = 10, 20$ against 1,023 (Montúfar et al. 2014; Telgarsky 2016).
- ▶ Depth pays when there is structure to fold: images, sound, text (Weeks 4 to 6). A 23-feature table has little to fold, so the sweep found one layer enough; deeper is also harder to train (Day 17).

Rule	Source	The catch
Start with one hidden layer, at least as wide as the input; give every hidden layer the same width. Make the network large enough to fit the training set, then control overfitting with weight decay, dropout and early stopping; shrink it last. If you can tune only one dial, tune the learning rate.	Bengio (2012), section 3.1 Goodfellow et al. (2016), section 11.4; Ng (2018), chapters 25 and 27 Goodfellow et al. (2016), section 11.4.1; Bengio (2012)	Wider than needed costs time, not much accuracy, as long as you stop early. Shrinking first brings underfitting back; regularisers keep the capacity and tame it. A bad rate looks like a wrong shape: a loss stuck at $\ln k$ or exploding (Day 17).
Try shapes by random search over a log-spaced range, not a full grid, when time is short. One hidden layer of 100 units is the default shape of scikit-learn's <code>MLPClassifier</code> .	Bergstra and Bengio (2012) scikit-learn documentation	Widths 32, 128, 512, not 100, 110, 120: the jumps must be big enough to see. A default for "any table", not a recommendation for yours.
Judge every shape on the validation score, never on the training loss or the test set.	Goodfellow et al. (2016), section 5.3	The hidden test set only ever sees the shape your validation split chose.

In the Olympiad solutions of our resources list, a recent table task (NEOAI 2025) was won with gradient-boosted trees, Week 2's tool. When a network is the right tool for a table, its shape is one or two hidden layers chosen by the loop on the next frame; on 3,000 rows a tower of layers only memorises faster, as the sweep showed.




1. **Baseline first:** inputs-128-outputs, ReLU, Adam at 0.001, early stopping on validation. Measured: 23-128-1 scores 6.40 at epoch 41 (3,201 dials), on the sweep's plateau.
2. **Read the two errors**, change one thing. Training error too high: more capacity, width before depth. Validation far above training: regularise and stop earlier; shrink last.
3. **Budget:** count the dials, time one epoch. A shape you cannot train three times before the deadline is too big.
4. **Keep the smallest shape within noise of the best:** here one layer of 16 to 64 units (6.39); 1,024 units or six layers only tie it at 14 to 16 times the dials. Trust a 0.1 difference only after a second seed.

The loop of Goodfellow et al. (2016), chapter 11: baseline, instrument, change one thing, measure.

Activations: **Output Layers and Their Losses**

The Double Softmax, Measured



Source: 784-128-10 on 20,000 MNIST images, Adam at 0.001, eight epochs; the only difference is an `nn.Softmax(dim=1)` as the model's last layer; `scripts/day18_activations_figures.py`.

- ▶ The double softmax trains: accuracy 0.957 after eight epochs against 0.964 for the right model. That is why it goes unnoticed. Its training loss, 1.49 against 0.035, is the tell: a probability vector pushed through a second softmax can never be confident, so the loss has a floor near 1.5.
- ▶ On a harder task the gap in accuracy grows and the model stops improving early. Rule: the last layer is `nn.Linear`; `CrossEntropyLoss` does the softmax.

The Pairing Table, Once More

Target	Last layer	Loss	To read a prediction
a number	Linear(h, 1), nothing after	L1Loss or MSELoss	the output itself
yes / no	Linear(h, 1), nothing after	BCEWithLogitsLoss	<code>torch.sigmoid(out)</code>
one of k	Linear(h, k), nothing after	CrossEntropyLoss	<code>softmax(out, 1)</code> or <code>argmax</code>

```
model = nn.Sequential(nn.Linear(23, 64), nn.ReLU(),  
                      nn.Linear(64, 1))      # a price: stop here  
pred = model(X)                             # dollars, any value  
probs = torch.sigmoid(model_yes_no(X))      # only to read a yes/no
```

- ▶ The rule has no exceptions in this course: **the last layer is linear**. The loss supplies the sigmoid or the softmax during training; you supply it, by hand, only to read probabilities.
- ▶ Yes/no needs one output, not two: a two-class softmax is the sigmoid of the logit difference, $e^{z_1} / (e^{z_1} + e^{z_2}) = \sigma(z_1 - z_2)$, so BCEWithLogitsLoss on one output trains the same model.
- ▶ The two classic mistakes: a softmax or sigmoid *inside* the model before the loss (the previous frame), and an activation on a regression output (the practice task).

Activations: **Olympiad Practice**

Task (Day 08 and Day 16, revisited)


The Steam prices, scored by MAE (full points below 6.6). A teammate trains a network for 100 epochs with Day 17's recipe and the L1 loss, and the validation MAE never drops below 15. Predicting the mean price for every game scores 8.49. Here is the model.

```
# "the sigmoid keeps the output positive"
model = nn.Sequential(nn.Linear(23, 64), nn.ReLU(),
                      nn.Linear(64, 1), nn.Sigmoid())
```

- ▶ **Question 1:** without running anything, what is the largest price this model can predict, and what does that do to a game priced at 60 dollars?
- ▶ **Question 2:** the teammate replaces `nn.Sigmoid()` by `nn.ReLU()` “because prices are never negative”. Does that fix it, and what new risk does it bring?
- ▶ **Question 3:** which single line gives the right model, and how do you check the fix from the predictions alone, before looking at the MAE?

Adapted from the ONIA (Romanian AI Olympiad) problem-proposal example “Predicting Game Prices on the Steam Platform”, 2024/25; data from the Kaggle Steam Games data

set.



Source: Day 16's features, 2,400 / 600 split, 23-64-1 networks, Adam at 0.001, L1 loss, 100 epochs; scripts/day18.activations_figures.py.

1. A sigmoid outputs numbers between 0 and 1, so the largest price it can predict is 1 dollar. Measured: every prediction is 1.00 and the MAE is 15.48, worse than the mean's 8.49. A 60-dollar game alone costs 59 dollars of error.
2. `nn.ReLU()` works here (MAE 6.43, predictions 2.80 to 55.48) because prices are positive, but it gains nothing over no activation (6.42) and adds a risk: if the output's z goes negative for every game, the output is dead at 0 and never recovers.
3. Delete the activation: `nn.Linear(64, 1)` is the last layer. Check `model(X_val).min()` and `.max()`: the predictions must span the real prices (0 to about 70 dollars), not a corner of them.

Activations: **Check Yourself**

Answer before turning the page.

1. **Definition.** A hidden unit's z is -30 for every row of the data. What is its output and its slope (a) with a sigmoid, (b) with ReLU, (c) with Leaky ReLU, and which of the three units can still recover during training?
2. **Mechanism.** A one-input ReLU network has hidden units $(w, b, v) = (1, -1, 2)$ and $(1, -3, -2)$ and output bias 0. Compute $\hat{y}$ at $x = 0, 1, 2, 3, 4$ and describe the shape of the curve.
3. **Code.** What is wrong with this model for a task scored by accuracy over 10 classes, and how does it show up during training?
4. **Design.** A 23-1024-1024-1 network reaches a training MAE of 2.27 and a validation MAE of 7.83 after 300 epochs; the 23-64-1 baseline had 5.49 and 6.42. Name the problem, the first change to try, and the shape you would hand in.

```
model = nn.Sequential(nn.Linear(784, 128), nn.Sigmoid(),  
                      nn.Linear(128, 10), nn.Softmax(dim=1))  
loss_fn = nn.CrossEntropyLoss()
```

1. (a) Sigmoid: output $\sigma(-30) \approx 10^{-13}$, slope $\approx 10^{-13}$: saturated, in practice frozen. (b) ReLU: output 0, slope 0: dead; no gradient reaches its weights, so it cannot recover. (c) Leaky ReLU: output -0.3 , slope 0.01: small, but not zero, so the weights still move and the unit can come back. Only (c) recovers on its own; (a) and (b) need the inputs or the rate fixed upstream.
2. Unit 1 is $2\text{ReLU}(x - 1)$: kink at 1, slope 2 to the right. Unit 2 is $-2\text{ReLU}(x - 3)$: kink at 3, slope -2 to the right. $\hat{y}$ at $x = 0, 1, 2, 3, 4$: 0, 0, 2, 4, 4. Flat until 1, rising with slope 2 until 3, flat again: the second hinge cancels the first past $x = 3$. Two units, two kinks, three straight pieces.
3. Two problems. The `Softmax` inside the model followed by `CrossEntropyLoss` is a double softmax: the loss cannot fall below about 1.5 and the accuracy climbs slowly (the measured frame). The sigmoid hidden layer trains, but more slowly than ReLU, and would stall in a deeper network. Fix: `nn.ReLU()` inside, no activation after the last `Linear`, and read probabilities with `softmax` only after training.
4. Overfitting: a gap of 5.6 dollars against 0.9. The training error is not the problem, so more capacity is not the answer. First change: early stopping (its best validation MAE, 6.38 at epoch 8, ties the baseline) and regularisation. Hand in the baseline: the same score with 1,601 dials instead of 1,075,201, and a model you can retrain three times before the deadline.

Activations: **Limitations & What Comes Next**

- ▶ **The hinge picture is exact only with one input.** With 784 inputs the kinks become hyperplanes and the pieces regions you cannot draw; the mechanism is the same.
- ▶ **ReLU is a default, not a law.** Leaky ReLU and GELU win on some tasks by a little; sigmoid and tanh still live inside recurrent networks (Week 6), where a bounded output is the point.
- ▶ **Width and depth were measured on two small data sets.** MNIST rewarded width, the Steam table nothing past 64 units, convolutions (Week 4) reward depth. The rules of thumb come from images and text: on your table, measure. Huge networks can break “bigger overfits” (double descent, Nakkiran et al. 2020).
- ▶ **Dead units are a symptom, not a disease.** The count says a rate or a scaling is wrong; it is not a dial.

The common mistake

Swapping activations or shapes at random. Read the dashboard first: a loss at $\ln k$ from the start is a rate or scaling problem, a loss near 1.5 on ten classes a double softmax, a gap between the two errors a capacity problem.

- ▶ **Day 19, self-study:** tensors, datasets and DataLoaders in depth: shapes, dtypes, custom datasets with transforms, and the loading speed that decides Week 4's image runs.
- ▶ **Day 20, self-study, Olympiad practice:** reading loss curves, the optimizer and regularisation dials, and the habits that win points on Friday.
- ▶ **Week 4:** the same neurons and the same activations, arranged as convolutions; the first thing you will check in every network is still its last layer and its loss.

Activations: **Summary**

- ▶ **Slopes:** sigmoid never above 0.25 and near 0 past $|z| = 4$, tanh the same past $|z| = 3$, ReLU exactly 1 or 0. Multiplied through five sigmoid layers, even the best case is 0.001: the vanishing gradient, measured as a 4,900-fold gap between the last and the first layer.
- ▶ **Saturated and dead units** come from unscaled inputs and large rates; the fixes are upstream (scaling, BatchNorm, a smaller rate), with Leaky ReLU as the activation that can recover.
- ▶ **A ReLU network is a sum of hinges:** one kink per hidden unit, straight pieces between; six units fit 3,598 used-car prices better than a line, and the picture shows which unit bends where.
- ▶ **Width and depth are design choices read off two errors:** training error high, add capacity (width first); validation far above training, regularise and stop early, shrink last; keep the smallest shape within noise of the best. Bengio, Goodfellow and Ng say the same, and the Steam sweep measured it.
- ▶ **The last layer is linear:** the loss supplies the sigmoid or softmax; a softmax inside the model or an activation on a regression output are the two mistakes to check first.

Activations: **References & Resources**

- [1] Glorot, X. and Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. *AISTATS*. Vanishing gradients through saturating activations.
- [2] Nair, V. and Hinton, G. (2010). Rectified linear units improve restricted Boltzmann machines. *ICML*. ReLU.
- [3] Maas, A., Hannun, A. and Ng, A. (2013). Rectifier nonlinearities improve neural network acoustic models. *ICML workshop*. Leaky ReLU.
- [4] Paszke, A. et al. (2019). PyTorch: an imperative style, high-performance deep learning library. *NeurIPS* 32. `torch.nn` activations and losses. <https://pytorch.org/docs>
- [5] LeCun, Y., Cortes, C. and Burges, C. *The MNIST database of handwritten digits*, via OpenML (`mnist_784`); the Day 15 split for every MNIST run in this deck.
- [6] Bin Talib, T. *Saudi Arabia Used Cars Dataset*, Kaggle `turkibintalib/saudi-arabia-used-cars-dataset`; the 3,598 listings from 2015 on with at most 300,000 km, for the hinge figure.
- [7] ONIA (Romanian AI Olympiad), official problem-proposal example *Predicting Game Prices on the Steam Platform*, Olimpiada-AI, 2024/25. The practice task.

- [8] Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*, MIT Press. Chapter 5 (capacity, underfitting, overfitting), chapter 6 (hidden and output units), chapter 11 (practical methodology: baseline, instrument, tune the learning rate first, add capacity then regularise).
<https://www.deeplearningbook.org>
- [9] Bengio, Y. (2012). Practical recommendations for gradient-based training of deep architectures. In *Neural Networks: Tricks of the Trade*, 2nd ed., Springer. Section 3.1: the number of hidden units, the same width per layer, the learning rate as the most important hyperparameter.
- [10] Ng, A. (2018). *Machine Learning Yearning*, deeplearning.ai. Chapters 25 and 27: techniques for reducing bias and variance.
- [11] Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems* 2; Hornik, K., Stinchcombe, M. and White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural Networks* 2. One hidden layer suffices in principle.
- [12] Montúfar, G., Pascanu, R., Cho, K. and Bengio, Y. (2014). On the number of linear regions of deep neural networks. *NeurIPS* 27; Telgarsky, M. (2016). Benefits of depth in neural networks. *COLT*. Depth multiplies the straight pieces (the folding frame).
- [13] Bergstra, J. and Bengio, Y. (2012). Random search for hyper-parameter optimization. *JMLR* 13.

- [14] Nakkiran, P., Kaplun, G., Bansal, Y., Yang, T., Barak, B. and Sutskever, I. (2020). Deep double descent: where bigger models and more data hurt. *ICLR*.
- [15] scikit-learn developers. `sklearn.neural_network.MLPClassifier`, default `hidden_layer_sizes=(100,)`.
https://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html
- [16] All plots and every number quoted: `slides-builder/scripts/day18_activations_figures.py` (gradients, hinges, shapes, dead units, double softmax, the practice task) and `day18_design_figures.py` (the three regimes, the width and depth sweep, the folding picture, the hinge network unit by unit).